Data Science with PySpark

David Kearney

Aug 22, 2020

CONTENTS

Data Science with PySpark, written by David R. Kearney.

Note: Data Science with PySpark

PYSPARK REGRESSION WITH FISCAL DATA

"A minimal example of using Pyspark for Linear Regression"

- toc: true- branch: master- badges: true

- comments: true

- author: David Kearney

- categories: [pyspark, jupyter]

- description: A minimal example of using Pyspark for Linear Regression

- title: Pyspark Regression with Fiscal Data

1.1 Bring in needed imports

```python
from pyspark.sql.functions import col
from pyspark.sql.types import StringType,BooleanType,DateType,IntegerType
from pyspark.sql.functions import *
```

1.2 Load data from CSV

```python
#collapse-hide

# Load data from a CSV
file_location = "/FileStore/tables/df_panel_fix.csv"
df = spark.read.format("CSV").option("inferSchema", True).option("header", True).
↪load(file_location)
display(df.take(5))
```

```python
df.createOrReplaceTempView("fiscal_stats")

sums = spark.sql("""
select year, sum(it) as total_yearly_it, sum(fr) as total_yearly_fr
from fiscal_stats
group by 1
order by year asc
""")

sums.show()
```

1.3 Describing the Data

```
df.describe().toPandas().transpose()
```

1.4 Cast Data Type

```
df2 = df.withColumn("gdp",col("gdp").cast(IntegerType())) \
.withColumn("specific",col("specific").cast(IntegerType())) \
.withColumn("general",col("general").cast(IntegerType())) \
.withColumn("year",col("year").cast(IntegerType())) \
.withColumn("fdi",col("fdi").cast(IntegerType())) \
.withColumn("rnr",col("rnr").cast(IntegerType())) \
.withColumn("rr",col("rr").cast(IntegerType())) \
.withColumn("i",col("i").cast(IntegerType())) \
.withColumn("fr",col("fr").cast(IntegerType()))
```

1.5 printSchema

```
df2.printSchema()
```

```python
from pyspark.ml.feature import VectorAssembler
from pyspark.ml.regression import LinearRegression

assembler = VectorAssembler(inputCols=['gdp', 'fdi'], outputCol="features")
train_df = assembler.transform(df2)
```

```
train_df.select("specific", "year").show()
```

1.6 Linear Regression in Pyspark

```python
lr = LinearRegression(featuresCol = 'features', labelCol='it')
lr_model = lr.fit(train_df)

trainingSummary = lr_model.summary
print("Coefficients: " + str(lr_model.coefficients))
print("RMSE: %f" % trainingSummary.rootMeanSquaredError)
print("R2: %f" % trainingSummary.r2)
```

```python
lr_predictions = lr_model.transform(train_df)
lr_predictions.select("prediction","it","features").show(5)
from pyspark.ml.evaluation import RegressionEvaluator
lr_evaluator = RegressionEvaluator(predictionCol="prediction", \
                 labelCol="it",metricName="r2")
```

```
print("R Squared (R2) on test data = %g" % lr_evaluator.evaluate(lr_predictions))
```

```python
print("numIterations: %d" % trainingSummary.totalIterations)
print("objectiveHistory: %s" % str(trainingSummary.objectiveHistory))
trainingSummary.residuals.show()
```

```python
predictions = lr_model.transform(test_df)
predictions.select("prediction","it","features").show()
```

```python
from pyspark.ml.regression import DecisionTreeRegressor
dt = DecisionTreeRegressor(featuresCol ='features', labelCol = 'it')
dt_model = dt.fit(train_df)
dt_predictions = dt_model.transform(train_df)
dt_evaluator = RegressionEvaluator(
    labelCol="it", predictionCol="prediction", metricName="rmse")
rmse = dt_evaluator.evaluate(dt_predictions)
print("Root Mean Squared Error (RMSE) on test data = %g" % rmse)
```

```python
from pyspark.ml.regression import GBTRegressor
gbt = GBTRegressor(featuresCol = 'features', labelCol = 'it', maxIter=10)
gbt_model = gbt.fit(train_df)
gbt_predictions = gbt_model.transform(train_df)
gbt_predictions.select('prediction', 'it', 'features').show(5)


gbt_evaluator = RegressionEvaluator(
    labelCol="it", predictionCol="prediction", metricName="rmse")
rmse = gbt_evaluator.evaluate(gbt_predictions)
print("Root Mean Squared Error (RMSE) on test data = %g" % rmse)
```

GROUP BY AND AGGREGATION WITH PYSPARK

"Group By and Aggregation with Pyspark"

- toc: true- branch: master- badges: true

- comments: true

- author: David Kearney

- categories: [pyspark, jupyter]

- description: Group By and Aggregation with Pyspark

- title: Group By and Aggregation with Pyspark

```python
from pyspark.sql import SparkSession
from pyspark.sql.functions import countDistinct, avg,stddev


# Load data from a CSV
file_location = "/FileStore/tables/df_panel_fix.csv"
df = spark.read.format("CSV").option("inferSchema", True).option("header", True).
↪load(file_location)
display(df.take(5))
```

```python
df.printSchema()
```

```python
df.groupBy("province")
```

```python
df.groupBy("province").mean().show()
```

```python
df.groupBy("reg").mean().show()
```

```python
# Count
df.groupBy("reg").count().show()
```

```python
# Max
df.groupBy("reg").max().show()
```

```python
# Min
df.groupBy("reg").min().show()
```

```python
# Sum
df.groupBy("reg").sum().show()
```

```
# Max it across everything
df.agg({'specific':'max'}).show()
```

```
grouped = df.groupBy("reg")
grouped.agg({"it":'max'}).show()
```

```
df.select(countDistinct("reg")).show()
```

```
df.select(countDistinct("reg").alias("Distinct Region")).show()
```

```
df.select(avg('specific')).show()
```

```
df.select(stddev("specific")).show()
```

```
from pyspark.sql.functions import format_number
```

```
specific_std = df.select(stddev("specific").alias('std'))
specific_std.show()
```

```
specific_std.select(format_number('std',0)).show()
```

```
df.orderBy("specific").show()
```

```
df.orderBy(df["specific"].desc()).show()
```

THREE

HANDLING MISSING DATA WITH PYSPARK

```python
from pyspark.sql import SparkSession
from pyspark.sql.functions import countDistinct, avg,stddev

# Load data from a CSV
file_location = "/FileStore/tables/df_panel_fix.csv"
df = spark.read.format("CSV").option("inferSchema", True).option("header", True).
↪load(file_location)
display(df.take(5))
```

```python
df.show()
```

```python
# Has to have at least 2 NON-null values
df.na.drop(thresh=2).show()
```

```python
# Drop any row that contains missing data
df.na.drop().show()
```

```python
df.na.drop(subset=["general"]).show()
```

```python
df.na.drop(how='any').show()
```

```python
df.na.drop(how='all').show()
```

```python
df.na.fill('example').show()
```

```python
df.na.fill(0).show()
```

```python
df.na.fill('example',subset=['fr']).show()
```

```python
df.na.fill(0,subset=['general']).show()
```

```python
# Mean Imputation
from pyspark.sql.functions import mean
mean_val = df.select(mean(df['general'])).collect()
```

```python
mean_val[0][0]
```

```python
mean_gen = mean_val[0][0]
```

```
df.na.fill(mean_gen,["general"]).show()
```

```
df.na.fill(df.select(mean(df['general'])).collect()[0][0],['general']).show()
```

DATAFRAME FILITERING AND OPERATIONS WITH PYSPARK

```python
from pyspark.sql import SparkSession

# Load data from a CSV
file_location = "/FileStore/tables/df_panel_fix.csv"
df = spark.read.format("CSV").option("inferSchema", True).option("header", True).
↪load(file_location)
display(df.take(5))
```

```python
df.filter("specific<10000").show()
```

```python
df.filter("specific<10000").select('province').show()
```

```python
df.filter("specific<10000").select(['province','year']).show()
```

```python
df.filter(df["specific"] < 10000).show()
```

```python
df.filter((df["specific"] < 55000) & (df['gdp'] > 200) ).show()
```

```python
df.filter((df["specific"] < 55000) | (df['gdp'] > 20000) ).show()
```

```python
df.filter((df["specific"] < 55000) & ~(df['gdp'] > 20000) ).show()
```

```python
df.filter(df["specific"] == 8964.0).show()
```

```python
df.filter(df["province"] == "Zhejiang").show()
```

```python
df.filter(df["specific"] == 8964.0).collect()
```

```python
result = df.filter(df["specific"] == 8964.0).collect()
```

```python
type(result[0])
```

```python
row = result[0]
```

```python
row.asDict()
```

```python
for item in result[0]:
    print(item)
```

FIVE

DATAFRAMES, FORMATTING, CASTING DATA TYPE AND CORRELATION WITH PYSPARK

```python
from pyspark.sql import SparkSession

# Load data from a CSV
file_location = "/FileStore/tables/df_panel_fix.csv"
df = spark.read.format("CSV").option("inferSchema", True).option("header", True).
↪load(file_location)
display(df.take(5))
```

```python
df.columns
```

```python
df.printSchema()
```

```python
# for row in df.head(5):
#     print(row)
#     print('\n')
```

```python
df.describe().show()
```

```python
df.describe().printSchema()
```

```python
from pyspark.sql.functions import format_number
```

```python
result = df.describe()
result.select(result['province']
,format_number(result['specific'].cast('float'),2).alias('specific')
,format_number(result['general'].cast('float'),2).alias('general')
,format_number(result['year'].cast('int'),2).alias('year'),format_number(result['gdp
↪'].cast('float'),2).alias('gdp')
,format_number(result['rnr'].cast('int'),2).alias('rnr'),format_number(result['rr'].
↪cast('float'),2).alias('rr')
,format_number(result['fdi'].cast('int'),2).alias('fdi'),format_number(result['it'].
↪cast('float'),2).alias('it')
,result['reg'].cast('string').alias('reg')
            ).show()
```

```python
df2 = df.withColumn("specific_gdp_ratio",df["specific"]/(df["gdp"]*100))#.show()
```

```python
df2.select('specific_gdp_ratio').show()
```

```
df.orderBy(df["specific"].asc()).head(1)[0][0]
```

```python
from pyspark.sql.functions import mean
df.select(mean("specific")).show()
```

```python
from pyspark.sql.functions import max,min
```

```
df.select(max("specific"),min("specific")).show()
```

```
df.filter("specific < 60000").count()
```

```
df.filter(df['specific'] < 60000).count()
```

```python
from pyspark.sql.functions import count
result = df.filter(df['specific'] < 60000)
result.select(count('specific')).show()
```

```
(df.filter(df["gdp"]>8000).count()*1.0/df.count())*100
```

```python
from pyspark.sql.functions import corr
df.select(corr("gdp","fdi")).show()
```

```python
from pyspark.sql.functions import year
#yeardf = df.withColumn("Year",year(df["year"]))
```

```
max_df = df.groupBy('year').max()
```

```
max_df.select('year','max(gdp)').show()
```

```python
from pyspark.sql.functions import month
```

```python
#df.select("year","avg(gdp)").orderBy('year').show()
```

RDDS AND SCHEMAS AND DATA TYPES WITH PYSPARK

```python
from pyspark.sql import SparkSession

# Load data from a CSV
file_location = "/FileStore/tables/df_panel_fix.csv"
df = spark.read.format("CSV").option("inferSchema", True).option("header", True).
↪load(file_location)
display(df.take(5))
```

```python
df.show()
```

```python
df.printSchema()
```

```python
df.columns
```

```python
df.describe()
```

```python
from pyspark.sql.types import StructField,StringType,IntegerType,StructType
```

```python
data_schema = [
StructField("_c0", IntegerType(), True)
,StructField("province", StringType(), True)
,StructField("specific", IntegerType(), True)
,StructField("general", IntegerType(), True)
,StructField("year", IntegerType(), True)
,StructField("gdp", IntegerType(), True)
,StructField("fdi", IntegerType(), True)
,StructField("rnr", IntegerType(), True)
,StructField("rr", IntegerType(), True)
,StructField("i", IntegerType(), True)
,StructField("fr", IntegerType(), True)
,StructField("reg", StringType(), True)
,StructField("it", IntegerType(), True)
]
```

```python
final_struc = StructType(fields=data_schema)
```

```python
df = spark.read.format("CSV").schema(final_struc).load(file_location)
```

```python
df.printSchema()
```

```
df.show()
```

```
df['fr']
```

```
type(df['fr'])
```

```
df.select('fr')
```

```
type(df.select('fr'))
```

```
df.select('fr').show()
```

```
df.head(2)
```

```
df.select(['reg','fr'])
```

```
df.select(['reg','fr']).show()
```

```
df.withColumn('fiscal_revenue',df['fr']).show()
```

```
df.show()
```

```
df.withColumnRenamed('fr','new_fiscal_revenue').show()
```

```
df.withColumn('double_fiscal_revenue',df['fr']*2).show()
```

```
df.withColumn('add_fiscal_revenue',df['fr']+1).show()
```

```
df.withColumn('half_fiscal_revenue',df['fr']/2).show()
```

```
df.withColumn('half_fr',df['fr']/2)
```

```
df.createOrReplaceTempView("economic_data")
```

```
sql_results = spark.sql("SELECT * FROM economic_data")
```

```
sql_results
```

```
sql_results.show()
```

```
spark.sql("SELECT * FROM economic_data WHERE fr=634562").show()
```

WINDOW FUNCTIONS AND PIVOT TABLES WITH PYSPARK

```python
from pyspark.sql import SparkSession
from pyspark.sql.types import StructField,StringType,IntegerType,StructType,
↪DoubleType, FloatType
from pyspark.sql.functions import *


data_schema = [
StructField("_c0", IntegerType(), True)
,StructField("province", StringType(), True)
,StructField("specific", DoubleType(), True)
,StructField("general", DoubleType(), True)
,StructField("year", IntegerType(), True)
,StructField("gdp", FloatType(), True)
,StructField("fdi", FloatType(), True)
,StructField("rnr", DoubleType(), True)
,StructField("rr", FloatType(), True)
,StructField("i", FloatType(), True)
,StructField("fr", IntegerType(), True)
,StructField("reg", StringType(), True)
,StructField("it", IntegerType(), True)
]

final_struc = StructType(fields=data_schema)

file_location = "/FileStore/tables/df_panel_fix.csv"
df = spark.read.format("CSV").schema(final_struc).option("header", True).load(file_
↪location)

#df.printSchema()

df.show()
```

```python
df.limit(10).toPandas()
```

```python
df = df.withColumnRenamed("reg","region")
```

```python
df.limit(10).toPandas()
```

```python
# df = df.toDF(*['year', 'region', 'province', 'gdp', 'fdi', 'specific', 'general',
↪'it', 'fr', 'rnr', 'rr', 'i', '_c0', 'specific_classification', 'provinceIndex',
↪'regionIndex'])
```

```python
df = df.select('year','region','province','gdp', 'fdi')
```

```python
df.sort("gdp").show()
```

```python
from pyspark.sql import functions as F
df.sort(F.desc("gdp")).show()
```

```python
from pyspark.sql.types import IntegerType, StringType, DoubleType
df = df.withColumn('gdp', F.col('gdp').cast(DoubleType()))
```

```python
df = df.withColumn('province', F.col('province').cast(StringType()))
```

```python
df.filter((df.gdp>10000) & (df.region=='East China')).show()
```

```python
from pyspark.sql import functions as F

df.groupBy(["region","province"]).agg(F.sum("gdp") ,F.max("gdp")).show()
```

```python
df.groupBy(["region","province"]).agg(F.sum("gdp").alias("SumGDP"),F.max("gdp").alias(
 ↪"MaxGDP")).show()
```

```python
df.groupBy(["region","province"]).agg(
    F.sum("gdp").alias("SumGDP"),\
    F.max("gdp").alias("MaxGDP")\
    ).show()
```

```python
df.limit(10).toPandas()
```

```python
casesWithNewConfirmed = cases.withColumn("NewConfirmed", 100 + F.col("confirmed"))
casesWithNewConfirmed.show()
```

```python
df = df.withColumn("Exp_GDP", F.exp("gdp"))
df.show()
```

Note: Window functions

```python
# Window functions

from pyspark.sql.window import Window
windowSpec = Window().partitionBy(['province']).orderBy(F.desc('gdp'))
df.withColumn("rank",F.rank().over(windowSpec)).show()
```

```python
from pyspark.sql.window import Window
windowSpec = Window().partitionBy(['province']).orderBy('year')
```

```python
dfWithLag = df.withColumn("lag_7",F.lag("gdp", 7).over(windowSpec))
```

```python
df.filter(df.year>'2000').show()
```

```python
from pyspark.sql.window import Window

windowSpec = Window().partitionBy(['province']).orderBy('year').rowsBetween(-6,0)
```

```
dfWithRoll = df.withColumn("roll_7_confirmed",F.mean("gdp").over(windowSpec))
```

```
dfWithRoll.filter(dfWithLag.year>'2001').show()
```

```
from pyspark.sql.window import Window
windowSpec = Window().partitionBy(['province']).orderBy('year').rowsBetween(Window.
→unboundedPreceding,Window.currentRow)
```

```
dfWithRoll = df.withColumn("cumulative_gdp",F.sum("gdp").over(windowSpec))
```

```
dfWithRoll.filter(dfWithLag.year>'1999').show()
```

Note: Pivot Dataframes

```
pivoted_df = df.groupBy('year').pivot('province') \
                    .agg(F.sum('gdp').alias('gdp') , F.sum('fdi').alias('fdi'))
pivoted_df.limit(10).toPandas()
```

```
pivoted_df.columns
```

```
newColnames = [x.replace("-","_") for x in pivoted_df.columns]
```

```
pivoted_df = pivoted_df.toDF(*newColnames)
```

```
expression = ""
cnt=0
for column in pivoted_df.columns:
    if column!='year':
        cnt +=1
        expression += f"'{column}' , {column},"

expression = f"stack({cnt}, {expression[:-1]}) as (Type,Value)"
```

```
unpivoted_df = pivoted_df.select('year',F.expr(expression))
unpivoted_df.show()
```

EIGHT

LINEAR REGRESSION AND RANDOM FOREST/GBT CLASSIFICATION WITH PYSPARK

```python
from pyspark.sql import SparkSession
from pyspark.sql.types import StructField,StringType,IntegerType,StructType,
↪DoubleType, FloatType
from pyspark.sql.functions import *

data_schema = [
StructField("_c0", IntegerType(), True)
,StructField("province", StringType(), True)
,StructField("specific", DoubleType(), True)
,StructField("general", DoubleType(), True)
,StructField("year", IntegerType(), True)
,StructField("gdp", FloatType(), True)
,StructField("fdi", FloatType(), True)
,StructField("rnr", DoubleType(), True)
,StructField("rr", FloatType(), True)
,StructField("i", FloatType(), True)
,StructField("fr", IntegerType(), True)
,StructField("reg", StringType(), True)
,StructField("it", IntegerType(), True)
]

final_struc = StructType(fields=data_schema)

file_location = "/FileStore/tables/df_panel_fix.csv"
df = spark.read.format("CSV").schema(final_struc).option("header", True).load(file_
↪location)

#df.printSchema()

df.show()
```

```python
df.groupBy('province').count().show()
```

```python
mean_val = df.select(mean(df['general'])).collect()
mean_val[0][0]
mean_gen = mean_val[0][0]
df = df.na.fill(mean_gen,["general"])
```

```python
mean_val = df.select(mean(df['specific'])).collect()
mean_val[0][0]
```

(continues on next page)

```
mean_gen = mean_val[0][0]
df = df.na.fill(mean_gen,["specific"])
```

```
mean_val = df.select(mean(df['rr'])).collect()
mean_val[0][0]
mean_gen = mean_val[0][0]
df = df.na.fill(mean_gen,["rr"])
```

```
mean_val = df.select(mean(df['fr'])).collect()
mean_val[0][0]
mean_gen = mean_val[0][0]
df = df.na.fill(mean_gen,["fr"])
```

```
mean_val = df.select(mean(df['rnr'])).collect()
mean_val[0][0]
mean_gen = mean_val[0][0]
df = df.na.fill(mean_gen,["rnr"])
```

```
mean_val = df.select(mean(df['i'])).collect()
mean_val[0][0]
mean_gen = mean_val[0][0]
df = df.na.fill(mean_gen,["i"])
```

```
from pyspark.sql.functions import *
df = df.withColumn('specific_classification',when(df.specific >= 583470.7303370787,1).
→otherwise(0))
```

```
from pyspark.ml.feature import StringIndexer
```

```
indexer = StringIndexer(inputCol="province", outputCol="provinceIndex")
df = indexer.fit(df).transform(df)
```

```
indexer = StringIndexer(inputCol="reg", outputCol="regionIndex")
df = indexer.fit(df).transform(df)
```

```
df.show()
```

```
from pyspark.ml.linalg import Vectors
from pyspark.ml.feature import VectorAssembler
```

```
df.columns
```

```
assembler = VectorAssembler(
 inputCols=[
 'provinceIndex',
# 'specific',
 'general',
 'year',
 'gdp',
 'fdi',
 #'rnr',
 #'rr',
```

```
 #'i',
 #'fr',
 'regionIndex',
 'it'
 ],
 outputCol="features")
```

```
output = assembler.transform(df)
```

```
final_data = output.select("features", "specific")
```

```
train_data,test_data = final_data.randomSplit([0.7,0.3])
```

```
from pyspark.ml.regression import LinearRegression
lr = LinearRegression(labelCol='specific')
```

```
lrModel = lr.fit(train_data)
```

```
print("Coefficients: {} Intercept: {}".format(lrModel.coefficients,lrModel.intercept))
```

```
test_results = lrModel.evaluate(test_data)
```

```
print("RMSE: {}".format(test_results.rootMeanSquaredError))
print("MSE: {}".format(test_results.meanSquaredError))
print("R2: {}".format(test_results.r2))
```

```
from pyspark.sql.functions import corr
```

```
df.select(corr('specific','gdp')).show()
```

```
from pyspark.ml.classification import DecisionTreeClassifier,GBTClassifier,
→RandomForestClassifier
from pyspark.ml import Pipeline
```

```
dtc = DecisionTreeClassifier(labelCol='specific_classification',featuresCol='features
→')
rfc = RandomForestClassifier(labelCol='specific_classification',featuresCol='features
→')
gbt = GBTClassifier(labelCol='specific_classification',featuresCol='features')
```

```
final_data = output.select("features", "specific_classification")
train_data,test_data = final_data.randomSplit([0.7,0.3])
```

```
rfc_model = rfc.fit(train_data)
gbt_model = gbt.fit(train_data)
dtc_model = dtc.fit(train_data)
```

```
dtc_predictions = dtc_model.transform(test_data)
rfc_predictions = rfc_model.transform(test_data)
gbt_predictions = gbt_model.transform(test_data)
```

```python
from pyspark.ml.evaluation import MulticlassClassificationEvaluator
```

```python
acc_evaluator = MulticlassClassificationEvaluator(labelCol="specific_classification",
→predictionCol="prediction", metricName="accuracy")
```

```python
dtc_acc = acc_evaluator.evaluate(dtc_predictions)
rfc_acc = acc_evaluator.evaluate(rfc_predictions)
gbt_acc = acc_evaluator.evaluate(gbt_predictions)
```

```python
print('-'*80)
print('Decision tree accuracy: {0:2.2f}%'.format(dtc_acc*100))
print('-'*80)
print('Random forest ensemble accuracy: {0:2.2f}%'.format(rfc_acc*100))
print('-'*80)
print('GBT accuracy: {0:2.2f}%'.format(gbt_acc*100))
print('-'*80)
```

```python
df.select(corr('specific_classification','fdi')).show()
```

```python
df.select(corr('specific_classification','gdp')).show()
```